

Ashutosh Kumar Jha

12340390

January 19, 2025

### **Solution 1**

To solve this problem, we analyze the first digit  $x$  when sampling uniformly from the numbers 1 to 1000.

#### **Step 1: Distribution of $x$**

The first digit  $x$  can range from 1 to 9. For a uniform distribution from 1 to 1000, the probability  $p_x$  of each digit  $x$  is determined by the proportion of numbers in the range that begin with  $x$ . Specifically:

$$p_x = \frac{\text{Number of integers with first digit } x}{1000}.$$

Case-by-case analysis: - For  $x = 1$ , the numbers are 1, 10 to 19, 100 to 199:

$$\text{Count} = 1 + 10 + 100 = 111.$$

Therefore,  $p_1 = \frac{111}{1000} = 0.111$ .

- For  $x = 2$ , the numbers are 2, 20 to 29, 200 to 299:

$$\text{Count} = 1 + 10 + 100 = 111.$$

Similarly,  $p_2 = 0.111$ .

- The same logic applies for  $x = 3, 4, \dots, 9$ , as all have the same structure.

Hence, the probabilities are:

$$p_1 = p_2 = \dots = p_9 = 0.111.$$

**Step 2: Verification of Probabilities**

To ensure the probabilities sum to 1:

$$\sum_{x=1}^9 p_x = 9 \cdot 0.111 = 1.$$

Thus, the probabilities are correctly normalized.

**Step 3: Mean of  $x$** 

The mean  $\mu$  of  $x$  is given by:

$$\mu = \sum_{x=1}^9 x \cdot p_x.$$

Substituting  $p_x = 0.111$  for all  $x$ :

$$\mu = \sum_{x=1}^9 x \cdot 0.111 = 0.111 \cdot (1 + 2 + 3 + \cdots + 9).$$

The sum of integers from 1 to 9 is:

$$1 + 2 + \cdots + 9 = \frac{9 \cdot (9 + 1)}{2} = 45.$$

Therefore:

$$\mu = 0.111 \cdot 45 = 5.0.$$

**Step 4: Variance of  $x$** 

The variance  $\sigma^2$  of  $x$  is given by:

$$\sigma^2 = \sum_{x=1}^9 p_x \cdot (x - \mu)^2.$$

Expanding  $(x - \mu)^2$ :

$$\sigma^2 = \sum_{x=1}^9 p_x \cdot x^2 - \mu^2.$$

First, compute  $\sum_{x=1}^9 p_x \cdot x^2$ :

$$\sum_{x=1}^9 p_x \cdot x^2 = 0.111 \cdot (1^2 + 2^2 + \cdots + 9^2).$$

The sum of squares of integers from 1 to 9 is:

$$1^2 + 2^2 + \cdots + 9^2 = \frac{9 \cdot (9 + 1) \cdot (2 \cdot 9 + 1)}{6} = \frac{9 \cdot 10 \cdot 19}{6} = 285.$$

Thus:

$$\sum_{x=1}^9 p_x \cdot x^2 = 0.111 \cdot 285 = 31.635.$$

Finally, compute the variance:

$$\sigma^2 = 31.635 - (5.0)^2 = 31.635 - 25.0 = 6.635.$$

## Final Results


The probabilities are:

$$p_1 = p_2 = \cdots = p_9 = 0.111.$$

The mean and variance of  $x$  are:

$$\mu = 5.0, \quad \sigma^2 = 6.635.$$

## Diagram Representation



## Solution 2

We are tasked with finding the variance of the convolution of  $n$  identical Gaussian distributions  $N(0, \sigma^2)$ .

### Step 1: Prelude

The convolution of two Gaussian distributions  $N(0, \sigma_1^2)$  and  $N(0, \sigma_2^2)$  is also Gaussian, with:

$$\text{Variance} = \sigma_1^2 + \sigma_2^2.$$

This property holds because the Fourier transform of a Gaussian is another Gaussian, and the convolution theorem simplifies the multiplication of transforms.

### Step 2: Generalization for $n$ Gaussians

If we convolve  $n$  identical Gaussians, each with variance  $\sigma^2$ , the total variance is the sum of variances:

$$\text{Total Variance} = \sigma^2 + \sigma^2 + \cdots + \sigma^2 \quad (n \text{ terms}).$$

Thus:

$$\text{Total Variance} = n \cdot \sigma^2.$$

### Step 3: Verification

Let us verify this result using induction: - **Base Case** ( $n = 1$ ): For a single Gaussian  $N(0, \sigma^2)$ , the variance is  $\sigma^2$ , which satisfies  $n \cdot \sigma^2 = \sigma^2$ .

- **Inductive Step**: Assume for  $n = k$ , the total variance is  $k \cdot \sigma^2$ . For  $n = k + 1$ : - Convolve  $k$  Gaussians with variance  $k \cdot \sigma^2$  and one Gaussian with variance  $\sigma^2$ . - The total variance becomes:

$$k \cdot \sigma^2 + \sigma^2 = (k + 1) \cdot \sigma^2.$$

Thus, by induction, the formula holds for all  $n$ .


### Final Result

The variance of the convolution of  $n$  identical Gaussians  $N(0, \sigma^2)$  is:

$$\text{Total Variance} = n \cdot \sigma^2.$$

### Diagram Representation

Below is a visual representation of the convolution process, illustrating how variances add:



### Solution 3

We aim to explain why the probability distribution  $P(x)$  for the sum of two random variables is the convolution of their individual probability distributions.

#### Step 1: Definition of Convolution

The convolution of two probability distributions  $p_1(x)$  and  $p_2(x)$  is defined as:

$$P(x) = (p_1 * p_2)(x) = \int_{-\infty}^{\infty} p_1(t)p_2(x - t) dt.$$

This operation calculates the probability  $P(x)$  of obtaining the value  $x$  as a sum of two independent random variables with distributions  $p_1(x)$  and  $p_2(x)$ .

#### Step 2: Explanation Using Dice Example

Consider two six-sided dice, each having a uniform distribution over the numbers 1 to 6. The probability of rolling any specific number  $k$  on a die is:

$$p_1(k) = \frac{1}{6}, \quad k = 1, 2, 3, 4, 5, 6.$$

To compute the probability distribution  $P(x)$  of the sum  $x$  when rolling two dice, we consider all pairs  $(a, b)$  such that  $a + b = x$ , where  $a$  is the result of the first die and  $b$  is the result of the second die:

$$P(x) = \sum_a p_1(a) \cdot p_2(x - a).$$

Since the dice are independent, this is equivalent to the convolution of  $p_1(x)$  and  $p_2(x)$ .

### Step 3: Computation for Two Dice


For two dice, the possible sums  $x$  range from 2 to 12. The probabilities for each sum are:

$$P(2) = \frac{1}{36}, \quad P(3) = \frac{2}{36}, \quad P(4) = \frac{3}{36}, \quad \dots, \quad P(12) = \frac{1}{36}.$$

These probabilities are obtained by counting the number of pairs  $(a, b)$  that result in each sum  $x$  and multiplying by the individual probabilities  $\frac{1}{6}$ .

### Step 4: Visualization

The following diagram illustrates the convolution process and the resulting probabilities for the sum of two dice:



### Final Result

The probability distribution  $P(x)$  for the sum of two random variables is their convolution:

$$P(x) = \sum_a p_1(a) \cdot p_2(x - a),$$

and for the specific case of rolling two dice, the resulting probabilities for sums  $x = 2, \dots, 12$  are:

$$P(x) = \left( \frac{1}{36}, \frac{2}{36}, \frac{3}{36}, \frac{4}{36}, \frac{5}{36}, \frac{6}{36}, \frac{5}{36}, \frac{4}{36}, \frac{3}{36}, \frac{2}{36}, \frac{1}{36} \right).$$

## Solution 4

We aim to prove the following identity for random samples  $0 < x_1 < x_2 < \dots < x_n$  with probabilities  $p_1, p_2, \dots, p_n$ :

$$\text{mean} = E[x] = \sum_{i=1}^n p_i x_i = \int_{t=0}^{\infty} (\text{Probability that } x > t) dt.$$

### Step 1: Expected Value as a Weighted Sum

By definition, the expected value  $E[x]$  is calculated as:

$$E[x] = \sum_{i=1}^n p_i x_i,$$

where  $p_i$  represents the probability associated with  $x_i$ , and the terms  $x_1, x_2, \dots, x_n$  are ordered such that  $0 < x_1 < x_2 < \dots < x_n$ .

### Step 2: Probability that $x > t$

The probability that  $x > t$  can be expressed as:

$$\text{Probability that } x > t = \begin{cases} 1, & \text{if } t < x_1, \\ 1 - p_1, & \text{if } x_1 \leq t < x_2, \\ 1 - p_1 - p_2, & \text{if } x_2 \leq t < x_3, \\ \vdots & \\ 0, & \text{if } t \geq x_n. \end{cases}$$

### Step 3: Integration Approach

The expected value  $E[x]$  can also be expressed as:

$$E[x] = \int_{t=0}^{\infty} (\text{Probability that } x > t) dt.$$

Substituting the piecewise definition of the probability, the integral can be broken into segments:

$$E[x] = \int_{t=0}^{x_1} 1 dt + \int_{t=x_1}^{x_2} (1 - p_1) dt + \int_{t=x_2}^{x_3} (1 - p_1 - p_2) dt + \dots + \int_{t=x_{n-1}}^{x_n} p_n dt.$$

#### Step 4: Simplifying the Integral

Each segment  $\int_{t=x_{i-1}}^{x_i}$  evaluates to:

$$\int_{t=x_{i-1}}^{x_i} (1 - p_1 - \cdots - p_{i-1}) dt = (x_i - x_{i-1})(1 - p_1 - \cdots - p_{i-1}),$$

where  $x_0 = 0$ .

Summing these contributions gives:

$$E[x] = x_1 + (x_2 - x_1)(1 - p_1) + (x_3 - x_2)(1 - p_1 - p_2) + \cdots + (x_n - x_{n-1})p_n.$$

#### Step 5: Verification of Equality

Rearranging the terms, it can be verified that this summation is equivalent to:


$$E[x] = \sum_{i=1}^n p_i x_i.$$

Thus, we have shown that:

$$E[x] = \sum_{i=1}^n p_i x_i = \int_{t=0}^{\infty} (\text{Probability that } x > t) dt.$$

#### Visualization of the Probability Function

The piecewise nature of the probability function can be visualized as follows:





## Solution 5

Consider the Gaussian mixture model defined as:

$$0.4\mathcal{N}\left(\begin{bmatrix} 10 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}\right) + 0.6\mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 8.4 & 2.0 \\ 2.0 & 1.7 \end{bmatrix}\right).$$

### a. Marginal Distributions for Each Dimension

For a multivariate normal distribution, the marginal distribution of each dimension is also normal, with parameters derived from the corresponding elements of the mean vector and covariance matrix.

**\*\*Component 1:\*\***

$$\mathcal{N}\left(\begin{bmatrix} 10 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}\right)$$

- Marginal distribution for  $x_1$ :  $\mathcal{N}(10, 1)$  - Marginal distribution for  $x_2$ :  $\mathcal{N}(2, 1)$

**\*\*Component 2:\*\***

$$\mathcal{N}\left(\begin{bmatrix} 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 8.4 & 2.0 \\ 2.0 & 1.7 \end{bmatrix}\right)$$

- Marginal distribution for  $x_1$ :  $\mathcal{N}(0, 8.4)$  - Marginal distribution for  $x_2$ :  $\mathcal{N}(0, 1.7)$

The overall marginal distributions are weighted sums of the individual component marginals.

**\*\*Marginal Distribution for  $x_1$ :**

$$0.4\mathcal{N}(10, 1) + 0.6\mathcal{N}(0, 8.4)$$

**\*\*Marginal Distribution for  $x_2$ :**

$$0.4\mathcal{N}(2, 1) + 0.6\mathcal{N}(0, 1.7)$$

### b. Mean, Mode, and Median for Each Marginal Distribution

**\*\*Marginal Distribution for  $x_1$ :**

- **\*\*Mean:\*\***

$$E[x_1] = 0.4 \times 10 + 0.6 \times 0 = 4$$

- **\*\*Variance:\*\***

$$\text{Var}(x_1) = 0.4 \times (1 + (10 - 4)^2) + 0.6 \times (8.4 + (0 - 4)^2) = 25.04$$

- **Standard Deviation:**

$$\sigma_{x_1} = \sqrt{25.04} \approx 5.004$$

- **Median:** For a Gaussian mixture, the median does not have a closed-form solution and typically requires numerical methods for estimation.

- **Mode:** The mode of a Gaussian mixture is not straightforward to compute analytically and usually requires numerical optimization techniques.

**Marginal Distribution for  $x_2$ :**

- **Mean:**

$$E[x_2] = 0.4 \times 2 + 0.6 \times 0 = 0.8$$

- **Variance:**

$$\text{Var}(x_2) = 0.4 \times (1 + (2 - 0.8)^2) + 0.6 \times (1.7 + (0 - 0.8)^2) = 2.008$$

- **Standard Deviation:**

$$\sigma_{x_2} = \sqrt{2.008} \approx 1.417$$

- **Median:** As with  $x_1$ , the median requires numerical estimation.

- **Mode:** Determining the mode necessitates numerical optimization.

### c. Mean and Mode for the Two-Dimensional Distribution

- **Mean Vector:**

$$E[\mathbf{x}] = 0.4 \times \begin{bmatrix} 10 \\ 2 \end{bmatrix} + 0.6 \times \begin{bmatrix} 0 \\ 0 \end{bmatrix} = \begin{bmatrix} 4 \\ 0.8 \end{bmatrix}$$

- **Mode:** The mode of a Gaussian mixture in multiple dimensions is complex to determine analytically and typically requires numerical optimization methods.

## Solution 6

To model the stochastic behavior of your compiler's success ( $x = 1$ ) or failure ( $x = 0$ ), we assume a Bernoulli distribution with parameter  $\mu$ , representing the probability of success:

$$p(x \mid \mu) = \mu^x (1 - \mu)^{1-x}, \quad x \in \{0, 1\}$$

**\*\*Conjugate Prior:\*\***

The conjugate prior for the Bernoulli distribution is the Beta distribution, parameterized by  $\alpha$  and  $\beta$ :

$$p(\mu) = \frac{\mu^{\alpha-1}(1-\mu)^{\beta-1}}{B(\alpha, \beta)}$$

where  $B(\alpha, \beta)$  is the Beta function, serving as a normalization constant.

**\*\*Posterior Distribution:\*\***

Given a set of  $N$  independent observations  $\{x_1, x_2, \dots, x_N\}$ , the likelihood function is:

$$p(\{x_i\} \mid \mu) = \prod_{i=1}^N \mu^{x_i} (1-\mu)^{1-x_i}$$

Combining the prior and the likelihood, the unnormalized posterior is:

$$p(\mu \mid \{x_i\}) \propto p(\{x_i\} \mid \mu) \times p(\mu)$$

Substituting the expressions for the likelihood and the prior:

$$p(\mu \mid \{x_i\}) \propto \left( \prod_{i=1}^N \mu^{x_i} (1-\mu)^{1-x_i} \right) \times \mu^{\alpha-1} (1-\mu)^{\beta-1}$$

Simplifying the exponents by defining  $S = \sum_{i=1}^N x_i$  (the total number of successes) and  $F = N - S$  (the total number of failures):

$$p(\mu \mid \{x_i\}) \propto \mu^{S+\alpha-1} (1-\mu)^{F+\beta-1}$$

This is the kernel of a Beta distribution with updated parameters:

$$\mu \mid \{x_i\} \sim \text{Beta}(S + \alpha, F + \beta)$$

**\*\*Interpretation:\*\***

The posterior distribution  $p(\mu \mid \{x_i\})$  is a Beta distribution with parameters:

$$\alpha' = S + \alpha = \sum_{i=1}^N x_i + \alpha$$

$$\beta' = F + \beta = N - \sum_{i=1}^N x_i + \beta$$

These updated parameters combine the prior information ( $\alpha$  and  $\beta$ ) with the observed data ( $S$  successes and  $F$  failures), providing a new estimate for the probability of success  $\mu$ .

## Solution 7

To determine the probability that the mango was picked from Bag 2, we apply Bayes' theorem. Let:

- $B_1$ : Event that the fruit is picked from Bag 1.
- $B_2$ : Event that the fruit is picked from Bag 2.
- $M$ : Event that the fruit picked is a mango.

We aim to find  $P(B_2 | M)$ , the probability that the fruit was picked from Bag 2 given that it is a mango.

**Bayes' Theorem:**

$$P(B_2 | M) = \frac{P(M | B_2) \cdot P(B_2)}{P(M)}$$

**Calculating the Components:**

1. **Prior Probabilities:**

-  $P(B_1)$ : Probability of picking from Bag 1, which occurs if the coin shows heads.

$$P(B_1) = 0.6$$

-  $P(B_2)$ : Probability of picking from Bag 2, which occurs if the coin shows tails.

$$P(B_2) = 0.4$$

2. **Likelihoods:**

-  $P(M | B_1)$ : Probability of picking a mango from Bag 1. Bag 1 contains 4 mangoes and 2 apples, totaling 6 fruits.

$$P(M | B_1) = \frac{4}{6} = \frac{2}{3}$$

-  $P(M | B_2)$ : Probability of picking a mango from Bag 2. Bag 2 contains 4 mangoes and 4 apples, totaling 8 fruits.

$$P(M | B_2) = \frac{4}{8} = \frac{1}{2}$$

3. \*\*Total Probability of Picking a Mango,  $P(M)$ :\*\*

Using the law of total probability:

$$P(M) = P(M | B_1) \cdot P(B_1) + P(M | B_2) \cdot P(B_2)$$

Substituting the known values:

$$P(M) = \left( \frac{2}{3} \times 0.6 \right) + \left( \frac{1}{2} \times 0.4 \right)$$

Calculating each term:

$$\frac{2}{3} \times 0.6 = \frac{2 \times 0.6}{3} = \frac{1.2}{3} = 0.4$$

$$\frac{1}{2} \times 0.4 = \frac{0.4}{2} = 0.2$$

Adding these:

$$P(M) = 0.4 + 0.2 = 0.6$$

\*\*Applying Bayes' Theorem:\*\*

$$P(B_2 | M) = \frac{P(M | B_2) \cdot P(B_2)}{P(M)} = \frac{\left( \frac{1}{2} \right) \times 0.4}{0.6}$$

Calculating the numerator:

$$\frac{1}{2} \times 0.4 = \frac{0.4}{2} = 0.2$$

Thus:

$$P(B_2 | M) = \frac{0.2}{0.6} = \frac{1}{3} \approx 0.333$$

**\*\*Conclusion:\*\***

The probability that the mango was picked from Bag 2 is  $\frac{1}{3}$  or approximately 33.3%.

## Solution 8

Consider the time series model:

$$\begin{aligned}x_{t+1} &= \mathbf{A}x_t + \omega_t, & \omega_t &\sim \mathcal{N}(0, Q) \\ y_t &= \mathbf{C}x_t + v_t, & v_t &\sim \mathcal{N}(0, R)\end{aligned}$$

where  $\omega_t$  and  $v_t$  are independent Gaussian noise variables. Additionally, assume that the initial state distribution is  $p(x_0) = \mathcal{N}(\mu_0, \Sigma_0)$ .

**\*\* (a) Form of  $p(x_0, x_1, \dots, x_T)$ : \*\***

The joint distribution  $p(x_0, x_1, \dots, x_T)$  represents the probability of a specific sequence of states from time  $t = 0$  to  $t = T$ . Given the model's structure, this joint distribution is a multivariate Gaussian distribution. This is because:

- The initial state  $x_0$  is Gaussian:  $x_0 \sim \mathcal{N}(\mu_0, \Sigma_0)$ .
- Each subsequent state  $x_{t+1}$  is a linear transformation of  $x_t$  plus Gaussian noise  $\omega_t$ . Since linear transformations and sums of Gaussian variables remain Gaussian,  $x_{t+1}$  is Gaussian if  $x_t$  is Gaussian.

By induction, all states  $x_t$  are Gaussian, and their joint distribution is a multivariate Gaussian. The mean vector and covariance matrix of this distribution can be determined recursively using the system dynamics and noise covariances.

**\*\* (b) State Estimation: \*\***

Assume that the posterior distribution of the state at time  $t$ , given all observations up to time  $T$ , is Gaussian:

$$p(x_t \mid y_1, \dots, y_T) = \mathcal{N}(\mu_t, \Sigma_t)$$

**\*\* (b.1) Predicting  $p(x_{t+1} \mid y_1, \dots, y_T)$ : \*\***

To predict the state at time  $t+1$  given observations up to time  $T$ , we use the system's dynamics:

$$x_{t+1} = \mathbf{A}x_t + \omega_t$$

Taking the expectation and covariance:

$$\begin{aligned}\mathbb{E}[x_{t+1} \mid y_1, \dots, y_T] &= \mathbf{A}\mathbb{E}[x_t \mid y_1, \dots, y_T] = \mathbf{A}\mu_t \\ \text{Cov}(x_{t+1} \mid y_1, \dots, y_T) &= \mathbf{A}\Sigma_t\mathbf{A}^\top + Q\end{aligned}$$

Therefore, the predictive distribution is:

$$p(x_{t+1} \mid y_1, \dots, y_T) = \mathcal{N}(\mathbf{A}\mu_t, \mathbf{A}\Sigma_t\mathbf{A}^\top + Q)$$

\*\* (b.2) Predicting  $y_{t+1}$  Given  $y_1, \dots, y_T$ :\*\*

The observation at time  $t + 1$  is:

$$y_{t+1} = \mathbf{C}x_{t+1} + v_{t+1}$$

Given  $y_1, \dots, y_T$ , the predicted observation  $y_{t+1}$  has mean and covariance:

$$\begin{aligned}\mathbb{E}[y_{t+1} \mid y_1, \dots, y_T] &= \mathbf{C}\mathbb{E}[x_{t+1} \mid y_1, \dots, y_T] = \mathbf{C}\mathbf{A}\mu_t \\ \text{Cov}(y_{t+1} \mid y_1, \dots, y_T) &= \mathbf{C}\text{Cov}(x_{t+1} \mid y_1, \dots, y_T)\mathbf{C}^\top + R \\ &= \mathbf{C}(\mathbf{A}\Sigma_t\mathbf{A}^\top + Q)\mathbf{C}^\top + R\end{aligned}$$

Thus, the predictive distribution for  $y_{t+1}$  is:

$$p(y_{t+1} \mid y_1, \dots, y_T) = \mathcal{N}(\mathbf{C}\mathbf{A}\mu_t, \mathbf{C}(\mathbf{A}\Sigma_t\mathbf{A}^\top + Q)\mathbf{C}^\top + R)$$

\*\* (b.3) Updating  $p(x_{t+1} \mid y_1, \dots, y_{t+1})$  After Observing  $y_{t+1} = \hat{y}$ :\*\*

Upon observing  $y_{t+1} = \hat{y}$ , we update our estimate of  $x_{t+1}$  using the Kalman filter update equations:

- \*\*Kalman Gain:\*\*

$$K_{t+1} = \Sigma_{t+1|t}\mathbf{C}^\top(\mathbf{C}\Sigma_{t+1|t}\mathbf{C}^\top + R)^{-1}$$

where  $\Sigma_{t+1|t} = \mathbf{A}\Sigma_t\mathbf{A}^\top + Q$  is the prior covariance of  $x_{t+1}$ .

- \*\*Updated Mean:\*\*

$$\mu_{t+1} = \mu_{t+1|t} + K_{t+1}(\hat{y} - \mathbf{C}\mu_{t+1|t})$$

with  $\mu_{t+1|t} = \mathbf{A}\mu_t$  being the prior mean of  $x_{t+1}$ .

- \*\*Updated Covariance:\*\*

$$\Sigma_{t+1} = (\mathbf{I} - K_{t+1}\mathbf{C})\Sigma_{t+1|t}$$

Therefore, the updated distribution is:

$$p(x_{t+1} \mid y_1, \dots, y_{t+1}) =$$

## Solution 9

### Manipulation of Gaussian Random Variables

Consider a Gaussian random variable  $\mathbf{x} \sim \mathcal{N}(\boldsymbol{\mu}_x, \boldsymbol{\Sigma}_x)$ , where  $\mathbf{x} \in \mathbb{R}^D$ . Furthermore, we have

$$\mathbf{y} = A\mathbf{x} + \mathbf{b} + \boldsymbol{\omega},$$

where  $\mathbf{y} \in \mathbb{R}^E$ ,  $A \in \mathbb{R}^{E \times D}$ ,  $\mathbf{b} \in \mathbb{R}^E$ , and  $\boldsymbol{\omega} \sim \mathcal{N}(\mathbf{0}, Q)$  is independent Gaussian noise. “Independent” implies that  $\mathbf{x}$  and  $\boldsymbol{\omega}$  are independent random variables and that  $Q$  is diagonal.

#### a. Likelihood $p(\mathbf{y} \mid \mathbf{x})$

Given  $\mathbf{x}$ , the distribution of  $\mathbf{y}$  is determined by the linear transformation and the additive Gaussian noise  $\boldsymbol{\omega}$ . Therefore, the likelihood  $p(\mathbf{y} \mid \mathbf{x})$  is Gaussian:

$$p(\mathbf{y} \mid \mathbf{x}) = \mathcal{N}(\mathbf{y} \mid A\mathbf{x} + \mathbf{b}, Q).$$

#### b. Distribution $p(\mathbf{y})$

To find the distribution of  $\mathbf{y}$ , we note that it is a linear transformation of  $\mathbf{x}$  with added Gaussian noise  $\boldsymbol{\omega}$ . Since both  $\mathbf{x}$  and  $\boldsymbol{\omega}$  are Gaussian and independent,  $\mathbf{y}$  is also Gaussian. The mean and covariance of  $\mathbf{y}$  are given by:

$$\begin{aligned}\boldsymbol{\mu}_y &= A\boldsymbol{\mu}_x + \mathbf{b}, \\ \boldsymbol{\Sigma}_y &= A\boldsymbol{\Sigma}_x A^\top + Q.\end{aligned}$$

This follows from the properties of linear transformations of Gaussian random variables [?].

#### c. Measurement Mapping

Given the transformation

$$\mathbf{z} = C\mathbf{y} + \mathbf{v},$$

where  $\mathbf{z} \in \mathbb{R}^F$ ,  $C \in \mathbb{R}^{F \times E}$ , and  $\mathbf{v} \sim \mathcal{N}(\mathbf{0}, R)$  is independent Gaussian noise:

##### i. Likelihood $p(\mathbf{z} \mid \mathbf{y})$

The likelihood  $p(\mathbf{z} \mid \mathbf{y})$  is Gaussian:

$$p(\mathbf{z} \mid \mathbf{y}) = \mathcal{N}(\mathbf{z} \mid C\mathbf{y}, R).$$



## ii. Distribution $p(\mathbf{z})$

Since  $\mathbf{z}$  is a linear transformation of  $\mathbf{y}$  with added Gaussian noise  $\mathbf{v}$ , and  $\mathbf{y}$  is Gaussian,  $\mathbf{z}$  is also Gaussian. The mean and covariance of  $\mathbf{z}$  are:

$$\boldsymbol{\mu}_z = C\boldsymbol{\mu}_y = C(A\boldsymbol{\mu}_x + \mathbf{b}),$$

$$\boldsymbol{\Sigma}_z = C\boldsymbol{\Sigma}_yC^\top + R = C(A\boldsymbol{\Sigma}_xA^\top + Q)C^\top + R.$$

## d. Posterior Distribution $p(\mathbf{x} \mid \hat{\mathbf{y}})$

Given a measurement  $\hat{\mathbf{y}}$ , we seek the posterior distribution  $p(\mathbf{x} \mid \hat{\mathbf{y}})$ . Since  $\mathbf{x}$  and  $\mathbf{y}$  are jointly Gaussian, the posterior distribution is also Gaussian. The joint distribution  $p(\mathbf{x}, \mathbf{y})$  is:

$$\begin{bmatrix} \mathbf{x} \\ \mathbf{y} \end{bmatrix} \sim \mathcal{N} \left( \begin{bmatrix} \boldsymbol{\mu}_x \\ \boldsymbol{\mu}_y \end{bmatrix}, \begin{bmatrix} \boldsymbol{\Sigma}_x & \boldsymbol{\Sigma}_{xy} \\ \boldsymbol{\Sigma}_{yx} & \boldsymbol{\Sigma}_y \end{bmatrix} \right),$$

where  $\boldsymbol{\Sigma}_{xy} = \boldsymbol{\Sigma}_xA^\top$  and  $\boldsymbol{\Sigma}_{yx} = A\boldsymbol{\Sigma}_x$ .

The conditional distribution  $p(\mathbf{x} \mid \hat{\mathbf{y}})$  is Gaussian with mean and covariance:

$$\boldsymbol{\mu}_{\mathbf{x}|\hat{\mathbf{y}}} = \boldsymbol{\mu}_x + \boldsymbol{\Sigma}_{xy}\boldsymbol{\Sigma}_y^{-1}(\hat{\mathbf{y}} - \boldsymbol{\mu}_y),$$

$$\boldsymbol{\Sigma}_{\mathbf{x}|\hat{\mathbf{y}}} = \boldsymbol{\Sigma}_x - \boldsymbol{\Sigma}_{xy}\boldsymbol{\Sigma}_y^{-1}\boldsymbol{\Sigma}_{yx}.$$

These results are derived from the properties of conditional distributions of jointly Gaussian random variables

# Solution 10

## Part 1: Methodology

We aim to compute the standardized variable:

$$X = \frac{A_N - \frac{1}{2}}{2\sqrt{N}}$$

where  $A_N$  is the average of  $N$  random 0-1 samples.

### Sample Mean ( $A_N$ )

- Each sample is a Bernoulli random variable with parameter  $p = 0.5$ .
- The mean of a Bernoulli random variable is  $p$ , and the variance is  $p(1 - p)$ .
- For  $p = 0.5$ , both mean and variance are 0.25.
- The sample mean  $A_N$  is an unbiased estimator of  $p$ .

### Standardization

- The standard deviation of the sample mean  $A_N$  is  $\sqrt{\frac{p(1-p)}{N}} = \frac{1}{2\sqrt{N}}$ .
- Standardizing  $A_N$  involves subtracting the mean and dividing by the standard deviation:

$$Z = \frac{A_N - \frac{1}{2}}{\frac{1}{2\sqrt{N}}} = 2\sqrt{N} \left( A_N - \frac{1}{2} \right)$$

- Comparing this with the given expression for  $X$ , we see:

$$X = \frac{A_N - \frac{1}{2}}{2\sqrt{N}} = \frac{1}{4N} \times Z$$

- Therefore,  $X$  is a scaled version of the standard normal variable  $Z$ .

### Expected Value of $X$

- Since  $Z$  follows a standard normal distribution with mean 0, the expected value of  $X$  is also 0.

### Variance of $X$

- The variance of  $Z$  is 1.
- The variance of  $X$  is:

$$\text{Var}(X) = \left( \frac{1}{4N} \right)^2 \times \text{Var}(Z) = \frac{1}{16N^2}$$

## Part 2: Python Code

```
1 import numpy as np
2
3 # Number of samples
4 N = 1_000_000
5
6 # Generate N random samples from a Bernoulli
  distribution with p=0.5
7 samples = np.random.binomial(n=1, p=0.5, size=N)
8
9 # Compute the sample mean
10 A_N = np.mean(samples)
11
12 # Compute the standardized variable X
13 X = (A_N - 0.5) / (2 * np.sqrt(N))
14
15 # Output the results
16 print(f"Sample Mean (A_N): {A_N}")
17 print(f"Standardized Variable (X): {X}")
```

## Solution 11

### Part 1: Methodology

#### 1. Download and Summarize the Dataset

- The dataset contains measurements for three penguin species: Adelie, Gentoo, and Chinstrap.
- Features include:
  - species: Penguin species.
  - island: Island in the Palmer Archipelago.
  - culmen\_length\_mm: Bill length (mm).
  - culmen\_depth\_mm: Bill depth (mm).
  - flipper\_length\_mm: Flipper length (mm).

- `body_mass_g`: Body mass (g).
- `sex`: Gender of the penguin.
- The dataset may contain missing values that need to be addressed.

## 2. Exploratory Data Analysis (EDA)

- Visualize distributions of numerical features.
- Analyze relationships between features using scatter plots and pair plots.
- Examine differences between species.
- Handle missing values appropriately.

## 3. Identify Clusters

- Apply clustering algorithms such as K-Means to identify natural groupings.
- Use dimensionality reduction techniques like PCA for visualization.
- Evaluate clustering performance using metrics like silhouette score.

## 4. Predict Species of a Given Penguin

- Consider supervised learning algorithms:
  - Logistic Regression
  - Decision Trees
  - Random Forests
  - Support Vector Machines
  - k-Nearest Neighbors
- Evaluate models using cross-validation and metrics such as accuracy, precision, recall, and F1-score.
- Select the best model based on performance metrics.

## Part 2: Python Code

```
1 # Import necessary libraries
2 import pandas as pd
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5 from sklearn.cluster import KMeans
6 from sklearn.decomposition import PCA
7 from sklearn.preprocessing import StandardScaler
8 from sklearn.metrics import silhouette_score
9 from sklearn.model_selection import train_test_split
10 from sklearn.ensemble import RandomForestClassifier
11 from sklearn.metrics import classification_report,
    confusion_matrix
12
13 # Load the dataset
14 url = 'https://github.com/mwaskom/seaborn-data/raw/
    master/penguins.csv'
15 penguins = pd.read_csv(url)
16
17 # Display the first few rows of the dataset
18 print(penguins.head())
19
20 # Summary statistics
21 print(penguins.describe())
22
23 # Check for missing values
24 print(penguins.isnull().sum())
25
26 # Drop rows with missing values for simplicity
27 penguins.dropna(inplace=True)
28
29 # Exploratory Data Analysis (EDA)
30 # Pair plot to visualize relationships between features
31 sns.pairplot(penguins, hue='species')
32 plt.show()
33
34 # Box plots to visualize distributions
35 plt.figure(figsize=(12, 6))
```

```

36 sns.boxplot(x='species', y='culmen_length_mm', data=
    penguins)
37 plt.title('Culmen Length by Species')
38 plt.show()
39
40 # Clustering
41 # Select numerical features for clustering
42 features = ['culmen_length_mm', 'culmen_depth_mm', '
    flipper_length_mm', 'body_mass_g']
43 X = penguins[features]
44
45 # Standardize the features
46 scaler = StandardScaler()
47 X_scaled = scaler.fit_transform(X)
48
49 # Apply K-Means clustering
50 kmeans = KMeans(n_clusters=3, random_state=42)
51 penguins['cluster'] = kmeans.fit_predict(X_scaled)
52
53 # Evaluate clustering performance
54 silhouette_avg = silhouette_score(X_scaled, penguins['
    cluster'])
55 print(f'Silhouette Score: {silhouette_avg}')
56
57 # Visualize clusters using PCA
58 pca = PCA(n_components=2)
59 principal_components = pca.fit_transform(X_scaled)
60 penguins['pca1'] = principal_components[:, 0]
61 penguins['pca2'] = principal_components[:, 1]
62
63 plt.figure(figsize=(10, 6))
64 sns.scatterplot(x='pca1', y='pca2', hue='cluster', data=
    penguins, palette='Set1')
65 plt.title('PCA of Penguin Clusters')
66 plt.show()
67
68 # Predicting Species
69 # Encode categorical variables
70 penguins['sex'] = penguins['sex'].map({'Male': 0, '
    Female': 1})

```

```

71 penguins = pd.get_dummies(penguins, columns=['island'],
    drop_first=True)
72
73 # Define feature matrix and target vector
74 X = penguins.drop(columns=['species', 'cluster', 'pca1',
    'pca2'])
75 y = penguins['species']
76
77 # Split the data into training and testing sets
78 X_train, X_test, y_train, y_test = train_test_split(X, y
    , test_size=0.2, random_state=42)
79
80 # Train a Random Forest Classifier
81 rf = RandomForestClassifier(n_estimators=100,
    random_state=42)
82 rf.fit(X_train, y_train)
83
84 # Make predictions
85 y_pred = rf.predict(X_test)
86
87 # Evaluate the model
88 print(confusion_matrix(y_test, y_pred))
89 print(classification_report(y_test, y_pred))

```

## Dataset Collection and Analysis

The following dataset have been used to perform the analysis: [Download Dataset Here](#)

The colab notebook containing the code is attached below for reference: [Colab Notebook](#)

```

1 from google.colab import drive
2
3 drive.mount("/content/drive")
4
5 import pandas as pd
6 import numpy as np
7 import seaborn as sns
8 from scipy.stats import zscore

```

```

9 import matplotlib.pyplot as plt
10
11 df = pd.read_csv("/content/drive/MyDrive/Course Work/Sem
    4/Data Analysis and Visualization/Homework 1/
    panic_attack_dataset.csv")
12
13 df.head()
14
15 print("Dataset Information")
16 df.info()
17
18 print("Summary Statistics")
19 df.describe(include="all")
20
21 print("Missing Values")
22 df.isnull().sum()
23
24 # Dataset Cleaning
25
26 # Handling missing values i.e. filling numerical columns
    with means and categorical columns with mode
27 for col in df.select_dtypes(include=[np.number]).columns
    :
28     df[col].fillna(df[col].mean(), inplace=True)
29
30 for col in df.select_dtypes(include=[object]).columns:
31     df[col].fillna(df[col].mode()[0], inplace=True)
32
33 df.drop_duplicates(inplace=True)
34
35 df.describe(include="all")
36
37 for col in df.select_dtypes(include=[object]).columns:
38     print(col, df[col].unique())
39
40 for col in df.select_dtypes(include=[object]).columns:
41     df[col] = df[col].str.lower().str.strip()
42
43 for col in df.select_dtypes(include=[object]).columns:
44     print(col, df[col].unique())

```



```

45
46 # Outlier detection and removal
47 z_scores = np.abs(zscore(df.select_dtypes(include=[np.
    number])))
48 outliers = (z_scores > 3).any(axis=1)
49 print(outliers.sum())
50
51 df = df[~outliers]
52
53 # Univariate Analysis
54 # Numerical columns
55 for col in df.select_dtypes(include=[np.number]).columns
    :
56     plt.figure(figsize=(8, 5))
57     sns.histplot(df[col], kde=True, bins=30)
58     plt.title(f"Distribution of {col}")
59     plt.show()
60
61 for col in df.select_dtypes(include=[object]).columns:
62     plt.figure(figsize=(8, 5))
63     sns.countplot(data=df, x=col, order=df[col].
        value_counts().index)
64     plt.title(f"Distribution of {col}")
65     plt.xticks(rotation=45)
66     plt.show()
67
68 # Bivariate Analysis
69 numeric_df = df.select_dtypes(include=[np.number])
70 plt.figure(figsize=(10, 8))
71 sns.heatmap(numeric_df.corr(), annot=True, cmap="
    coolwarm", fmt=".2f")
72 plt.title("Correlation Matrix")
73 plt.show()
74
75 sns.pairplot(df.select_dtypes(include=[np.number]),
    diag_kind="kde")
76 plt.show()
77
78 numeric_df = df.select_dtypes(include=[np.number])
79 correlation_matrix = numeric_df.corr().unstack().

```

```
    sort_values(ascending=False)
80 correlation_matrix = correlation_matrix[
    correlation_matrix < 1]
81 print("Top Correlations: ")
82 print(correlation_matrix.head(10))
```