

# kumar-jha-12340390-dav-homework-2

February 5, 2025

## 0.1 Question 1

```
[2]: import pandas as pd
from google.colab import drive
import numpy as np
import matplotlib.pyplot as plt

drive.mount("/content/drive")

df = pd.read_csv("/content/drive/MyDrive/Course Work/Sem 4/Data Analysis and_
↳Visualization/Homework 2/DistanceTimeDataset - StudentsHomeTownDistance.csv")

cols = df.columns.tolist()
columns_to_drop = cols[6:]
df.drop(columns_to_drop, axis=1, inplace=True)

nan_df = df[df.isna().any(axis=1)]
rows_to_drop = nan_df.index.tolist()
df.drop(rows_to_drop, axis=0, inplace=True)
df.reset_index(drop=True, inplace=True)

## Normalizing the features
time = df[cols[1]]
distance = df[cols[2]]
time_normalized = (time - time.mean()) / time.std()
distance_normalized = (distance - distance.mean()) / distance.std()

X = time_normalized.values.reshape(-1, 1)
y = distance_normalized.values.reshape(-1, 1)

## Parameters for batch gradient descent
W = np.array([0, 35], dtype=float)
learning_rate = 0.01
num_epochs = 100
batch_size = 8
num_samples = X.shape[0]

gradient_history = []
```

```

loss_history = []
W0_history = []
W1_history = []
gradient_norm_history = []

for epoch in range(num_epochs):
    batch = np.random.choice(df.index, size=8, replace=True)
    X_batch = X[batch]
    y_batch = y[batch]
    y_pred = W[0] + W[1] * X_batch
    loss = np.mean((y_batch - y_pred) ** 2)
    loss_history.append(loss)

    gradient_W0 = -2 * np.sum(y_batch - y_pred) / batch_size
    gradient_W1 = -2 * np.sum((y_batch - y_pred) * X_batch) / batch_size
    gradient = np.array([gradient_W0, gradient_W1])
    gradient_history.append(gradient)

    W = W - learning_rate * gradient
    W0_history.append(W[0])
    W1_history.append(W[1])
    gradient_norm_history.append(np.linalg.norm(gradient))

gradient_history = np.array(gradient_history)

mean_gradient_W0 = np.mean(gradient_history[:, 0])
mean_gradient_W1 = np.mean(gradient_history[:, 1])
mean_gradient = np.array([mean_gradient_W0, mean_gradient_W1])

cov_matrix = np.cov(gradient_history.T)

y_pred_full = W[0] + W[1] * X
true_gradient_W0 = -2 * np.sum(y - y_pred_full) / num_samples
true_gradient_W1 = -2 * np.sum((y - y_pred_full) * X) / num_samples
true_gradient = np.array([true_gradient_W0, true_gradient_W1])

plt.figure(figsize=(10, 5))
plt.plot(gradient_history[:, 0], label="Gradient W0", linestyle='dashed')
plt.plot(gradient_history[:, 1], label='Gradient W1', linestyle='dashed')
plt.axhline(true_gradient[0], color='r', linestyle='solid', label='True_
↪Gradient W0')
plt.axhline(true_gradient[1], color='b', linestyle='solid', label='True_
↪Gradient W1')
plt.xlabel('Epochs')
plt.ylabel('Gradient Value')
plt.legend()
plt.title('Comparison of Batch Gradient Values and True Gradient Value')

```

```

plt.show()

# Plot Loss vs. Epochs
plt.figure()
plt.plot(loss_history, label='Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.title('Loss vs. Epochs')
plt.legend()
plt.show()

# Plot Loss vs. W0
plt.figure()
plt.plot(W0_history, loss_history, label='Loss vs. W0', marker='o')
plt.xlabel('W0')
plt.ylabel('Loss')
plt.title('Loss vs. W0 Path')
plt.legend()
plt.show()

# Plot Loss vs. W1
plt.figure()
plt.plot(W1_history, loss_history, label='Loss vs. W1', marker='o')
plt.xlabel('W1')
plt.ylabel('Loss')
plt.title('Loss vs. W1 Path')
plt.legend()
plt.show()

# Plot W0 vs. Epochs
plt.figure()
plt.plot(W0_history, label='W0 Value')
plt.xlabel('Epochs')
plt.ylabel('W0')
plt.title('W0 Value Change Over Epochs')
plt.legend()
plt.show()

# Plot W1 vs. Epochs
plt.figure()
plt.plot(W1_history, label='W1 Value')
plt.xlabel('Epochs')
plt.ylabel('W1')
plt.title('W1 Value Change Over Epochs')
plt.legend()
plt.show()


```


```

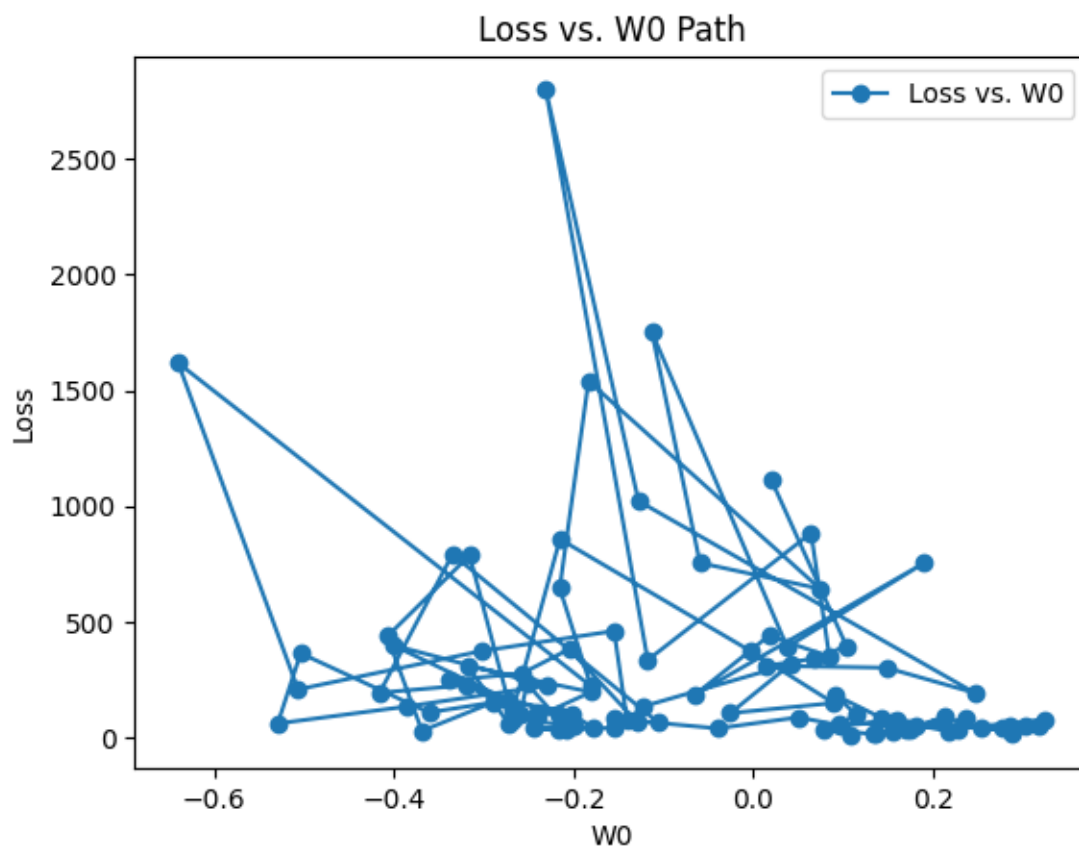
# Plot Loss vs. Gradient Norm
plt.figure()
plt.plot(gradient_norm_history, loss_history, label='Loss vs. Gradient Norm')
plt.xlabel('Gradient Norm')
plt.ylabel('Loss')
plt.title('Loss vs. Gradient Norm')
plt.legend()
plt.show()

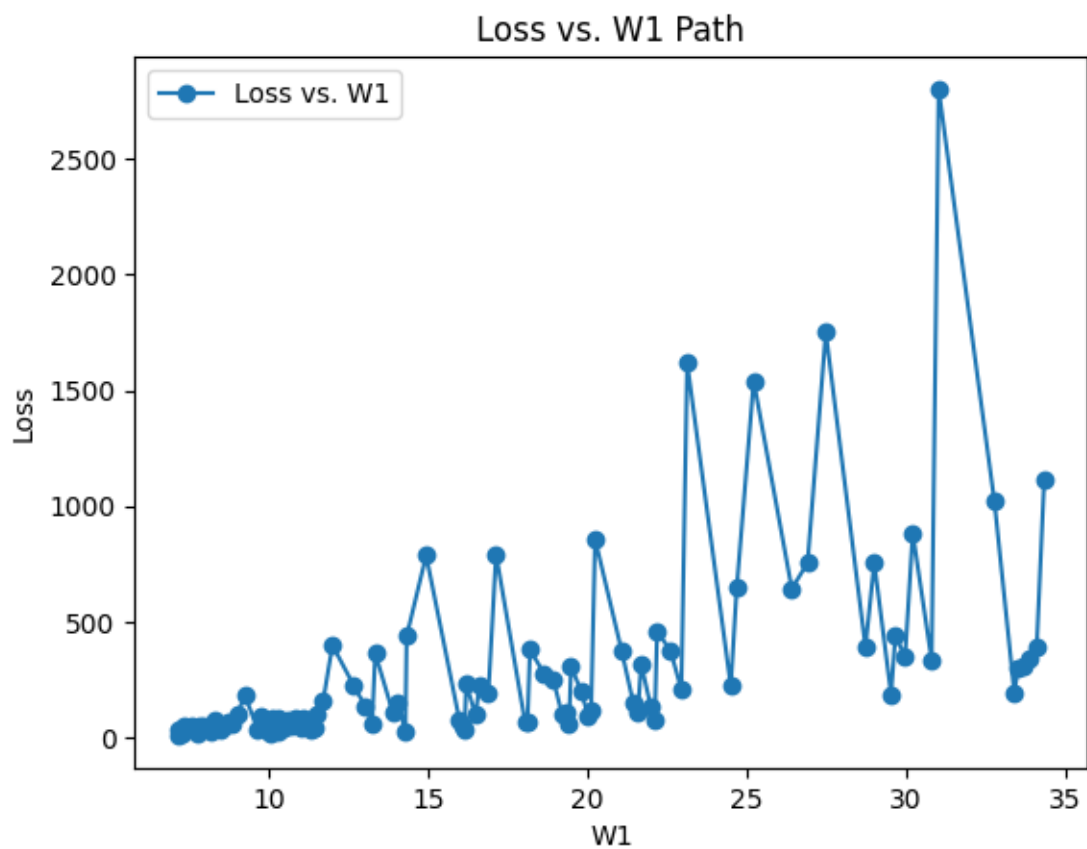
```


Drive already mounted at /content/drive; to attempt to forcibly remount, call `drive.mount("/content/drive", force_remount=True)`.






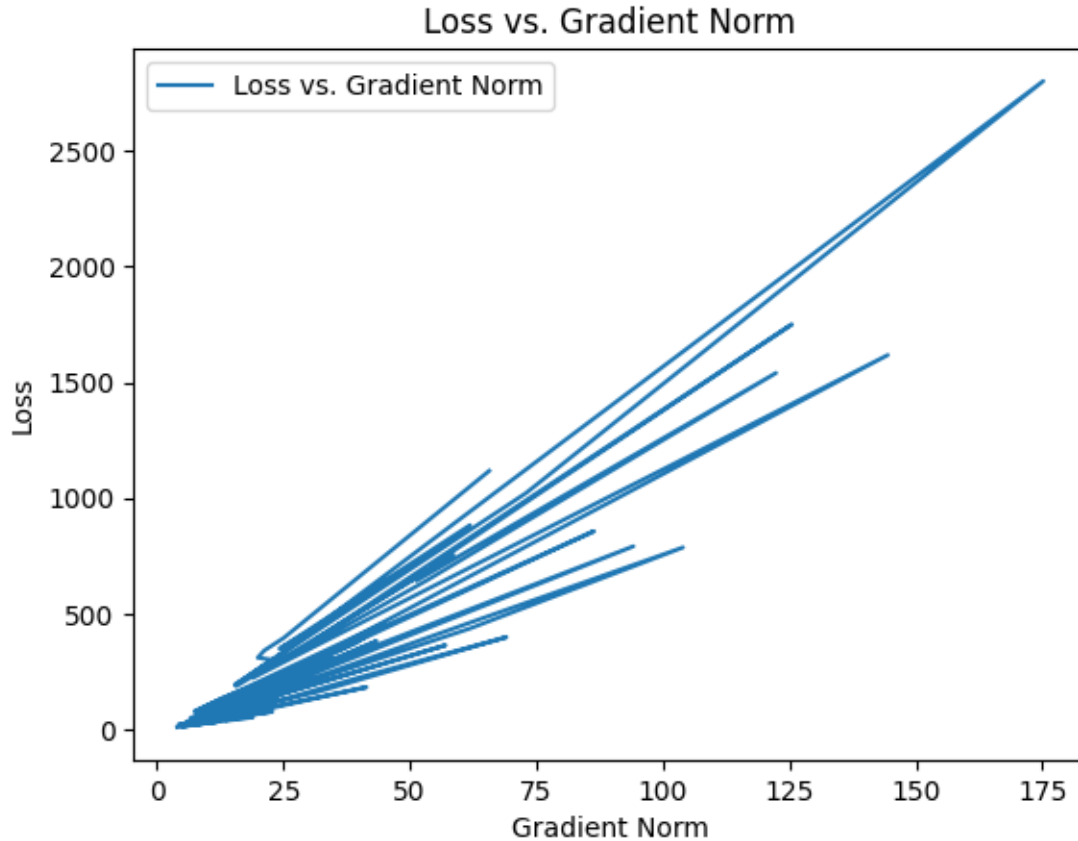












1. Random batches are created for 100 number of batches
2. Each random batch has a size of 8
3. Mean is calculated for both the gradients (gradients of both the parameters involved)
4. Covariance matrix is calculated keeping in account both the gradients
5. Visualization showing the true gradient values and also the gradient values vs number of epochs, for both the parameters involved
6. The path for both the gradient values are shown in the graph plotted
7. The path for both the parameter values are shown in the graph plotted, in terms of the number of epochs, and also in terms of the loss function
8. Variation of the loss function with the number of epochs is shown
9. The plot of the loss function with the gradient norm is shown
10. Answer to all the questions asked are present in the latex file submitted

## 1 Question 2

## 2 Submitted in latex, because those questions were Math Based

### 2.1 Question 3

```
[3]: import pandas as pd
      from google.colab import drive
      import numpy as np
      import matplotlib.pyplot as plt
      from sklearn.linear_model import LinearRegression
      from sklearn.metrics import mean_squared_error
      from sklearn.model_selection import train_test_split
```

```
[4]: drive.mount("/content/drive")
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True).

```
[5]: df = pd.read_csv("/content/drive/MyDrive/Course Work/Sem 4/Data Analysis and_
      ↪Visualization/Homework 2/DSAI CourseInterestRelevanceSurvey - Original.csv")
```

```
[6]: df.head()
```

```
[6]: Unnamed: 0  MAL100  MAL101  MAL403  EEL101  ECL101  BML101  CSL100  CSL201  \
0  Student 1      4.0      3.0      4.0      1.0      1.0      1.0      4.0      5.0
1  Student 2      3.0      3.0      3.0      1.0      1.0      1.0      4.0      5.0
2  Student 3      4.0      4.0      3.0      3.0      4.0      2.0      4.0      5.0
3  Student 4      3.0      4.0      4.0      1.0      1.0      1.0      5.0      5.0
4  Student 5      3.0      3.0      4.0      3.0      3.0      2.0      4.0      5.0
```

```
      CSL202  DSL201  Unnamed: 11  Unnamed: 12 Please fill values {1,2,3,4,5}  \
0      5.0      5.0              3.3              NaN                      Legend
1      5.0      5.0              3.1              NaN                      1
2      4.0      NaN              3.7              NaN                      2
3      5.0      2.0              3.1              NaN                      3
4      5.0      4.0              3.6              NaN                      4
```

```
      Unnamed: 14
0              NaN
1      Neither interesting nor relevant
2  Little interesting, but relevance not clear. 0...
3      Somewhat interesting and somewhat relevant
4              Interesting and Relevant
```

```
[7]: cols = df.columns.tolist()
```

```
[8]: cols
```

```
[8]: ['Unnamed: 0',
      'MAL100',
      'MAL101',
      'MAL403',
      'EEL101',
      'ECL101',
      'BML101',
      'CSL100',
      'CSL201',
      'CSL202',
      'DSL201',
      'Unnamed: 11',
      'Unnamed: 12',
      'Please fill values {1,2,3,4,5}',
      'Unnamed: 14']
```

```
[9]: df.rename(columns={cols[0]: "Student ID"}, inplace=True)
```

```
[10]: columns_to_drop = cols[12:]
```

```
[11]: df.drop(columns_to_drop, axis=1, inplace=True)
```

```
[12]: df.head()
```

```
[12]:
```

	Student ID	MAL100	MAL101	MAL403	EEL101	ECL101	BML101	CSL100	CSL201	\
0	Student 1	4.0	3.0	4.0	1.0	1.0	1.0	4.0	5.0	
1	Student 2	3.0	3.0	3.0	1.0	1.0	1.0	4.0	5.0	
2	Student 3	4.0	4.0	3.0	3.0	4.0	2.0	4.0	5.0	
3	Student 4	3.0	4.0	4.0	1.0	1.0	1.0	5.0	5.0	
4	Student 5	3.0	3.0	4.0	3.0	3.0	2.0	4.0	5.0	

	CSL202	DSL201	Unnamed: 11
0	5.0	5.0	3.3
1	5.0	5.0	3.1
2	4.0	NaN	3.7
3	5.0	2.0	3.1
4	5.0	4.0	3.6

```
[13]: df
```

```
[13]:
```

	Student ID	MAL100	MAL101	MAL403	EEL101	ECL101	BML101	CSL100	\
0	Student 1	4.0	3.0	4.0	1.0	1.0	1.0	4.0	
1	Student 2	3.0	3.0	3.0	1.0	1.0	1.0	4.0	
2	Student 3	4.0	4.0	3.0	3.0	4.0	2.0	4.0	
3	Student 4	3.0	4.0	4.0	1.0	1.0	1.0	5.0	
4	Student 5	3.0	3.0	4.0	3.0	3.0	2.0	4.0	
5	Student 6	3.0	3.0	5.0	1.0	1.0	1.0	4.0	

6	Student 7	3.0	4.0	4.0	1.0	1.0	1.0	4.0
7	Student 8	3.0	4.0	2.0	2.0	3.0	2.0	4.0
8	Student 9	3.0	4.0	4.0	2.0	3.0	1.0	4.0
9	Student 10	4.0	4.0	3.0	3.0	2.0	2.0	4.0
10	Student 11	4.0	5.0	4.0	2.0	3.0	3.0	5.0
11	Student 12	3.0	3.0	3.0	1.0	1.0	2.0	5.0
12	Student 13	2.0	3.0	3.0	1.0	1.0	2.0	4.0
13	Student 14	3.0	3.0	4.0	2.0	2.0	2.0	4.0
14	Student 15	4.0	4.0	5.0	3.0	3.0	1.0	4.0
15	Student 16	3.0	3.0	4.0	1.0	1.0	1.0	4.0
16	Student 17	4.0	4.0	3.0	1.0	1.0	1.0	5.0
17	Student 18	3.0	4.0	3.0	2.0	1.0	3.0	4.0
18	Student 19	3.0	3.0	4.0	1.0	2.0	1.0	4.0
19	Student 20	1.0	5.0	3.0	1.0	1.0	2.0	1.0
20	Student 21	3.0	4.0	3.0	2.0	3.0	1.0	4.0
21	Student 22	4.0	5.0	5.0	1.0	1.0	2.0	4.0
22	Student 23	2.0	2.0	2.0	1.0	1.0	4.0	4.0
23	Student 24	3.0	5.0	5.0	2.0	1.0	1.0	4.0
24	NaN	3.1	3.7	3.6	1.6	1.8	1.7	4.0

	CSL201	CSL202	DSL201	Unnamed: 11
0	5.0	5.0	5.0	3.3
1	5.0	5.0	5.0	3.1
2	5.0	4.0	NaN	3.7
3	5.0	5.0	2.0	3.1
4	5.0	5.0	4.0	3.6
5	5.0	5.0	5.0	3.3
6	4.0	4.0	5.0	3.1
7	4.0	4.0	5.0	3.3
8	4.0	4.0	4.0	3.3
9	4.0	5.0	4.0	3.5
10	5.0	5.0	4.0	4.0
11	4.0	3.0	5.0	3.0
12	4.0	3.0	3.0	2.6
13	5.0	4.0	4.0	3.3
14	5.0	4.0	5.0	3.8
15	4.0	4.0	4.0	2.9
16	5.0	5.0	5.0	3.4
17	4.0	4.0	5.0	3.3
18	3.0	5.0	5.0	3.1
19	5.0	4.0	3.0	2.6
20	4.0	4.0	5.0	3.3
21	1.0	5.0	3.0	3.1
22	2.0	2.0	4.0	2.4
23	4.0	4.0	5.0	3.4
24	4.2	4.3	4.3	3.2

```
[14]: df.drop([24],axis=0,inplace=True)
```

```
[15]: df
```

```
[15]:
```

	Student ID	MAL100	MAL101	MAL403	EEL101	ECL101	BML101	CSL100	\
0	Student 1	4.0	3.0	4.0	1.0	1.0	1.0	4.0	
1	Student 2	3.0	3.0	3.0	1.0	1.0	1.0	4.0	
2	Student 3	4.0	4.0	3.0	3.0	4.0	2.0	4.0	
3	Student 4	3.0	4.0	4.0	1.0	1.0	1.0	5.0	
4	Student 5	3.0	3.0	4.0	3.0	3.0	2.0	4.0	
5	Student 6	3.0	3.0	5.0	1.0	1.0	1.0	4.0	
6	Student 7	3.0	4.0	4.0	1.0	1.0	1.0	4.0	
7	Student 8	3.0	4.0	2.0	2.0	3.0	2.0	4.0	
8	Student 9	3.0	4.0	4.0	2.0	3.0	1.0	4.0	
9	Student 10	4.0	4.0	3.0	3.0	2.0	2.0	4.0	
10	Student 11	4.0	5.0	4.0	2.0	3.0	3.0	5.0	
11	Student 12	3.0	3.0	3.0	1.0	1.0	2.0	5.0	
12	Student 13	2.0	3.0	3.0	1.0	1.0	2.0	4.0	
13	Student 14	3.0	3.0	4.0	2.0	2.0	2.0	4.0	
14	Student 15	4.0	4.0	5.0	3.0	3.0	1.0	4.0	
15	Student 16	3.0	3.0	4.0	1.0	1.0	1.0	4.0	
16	Student 17	4.0	4.0	3.0	1.0	1.0	1.0	5.0	
17	Student 18	3.0	4.0	3.0	2.0	1.0	3.0	4.0	
18	Student 19	3.0	3.0	4.0	1.0	2.0	1.0	4.0	
19	Student 20	1.0	5.0	3.0	1.0	1.0	2.0	1.0	
20	Student 21	3.0	4.0	3.0	2.0	3.0	1.0	4.0	
21	Student 22	4.0	5.0	5.0	1.0	1.0	2.0	4.0	
22	Student 23	2.0	2.0	2.0	1.0	1.0	4.0	4.0	
23	Student 24	3.0	5.0	5.0	2.0	1.0	1.0	4.0	

	CSL201	CSL202	DSL201	Unnamed: 11
0	5.0	5.0	5.0	3.3
1	5.0	5.0	5.0	3.1
2	5.0	4.0	NaN	3.7
3	5.0	5.0	2.0	3.1
4	5.0	5.0	4.0	3.6
5	5.0	5.0	5.0	3.3
6	4.0	4.0	5.0	3.1
7	4.0	4.0	5.0	3.3
8	4.0	4.0	4.0	3.3
9	4.0	5.0	4.0	3.5
10	5.0	5.0	4.0	4.0
11	4.0	3.0	5.0	3.0
12	4.0	3.0	3.0	2.6
13	5.0	4.0	4.0	3.3
14	5.0	4.0	5.0	3.8
15	4.0	4.0	4.0	2.9

16	5.0	5.0	5.0	3.4
17	4.0	4.0	5.0	3.3
18	3.0	5.0	5.0	3.1
19	5.0	4.0	3.0	2.6
20	4.0	4.0	5.0	3.3
21	1.0	5.0	3.0	3.1
22	2.0	2.0	4.0	2.4
23	4.0	4.0	5.0	3.4

```
[16]: cols = df.columns.tolist()
n = len(cols)
column_to_drop = cols[n-1]
df.drop([column_to_drop],axis=1,inplace=True)
```

```
[17]: df
```

```
[17]:
```

	Student ID	MAL100	MAL101	MAL403	EEL101	ECL101	BML101	CSL100	\
0	Student 1	4.0	3.0	4.0	1.0	1.0	1.0	4.0	
1	Student 2	3.0	3.0	3.0	1.0	1.0	1.0	4.0	
2	Student 3	4.0	4.0	3.0	3.0	4.0	2.0	4.0	
3	Student 4	3.0	4.0	4.0	1.0	1.0	1.0	5.0	
4	Student 5	3.0	3.0	4.0	3.0	3.0	2.0	4.0	
5	Student 6	3.0	3.0	5.0	1.0	1.0	1.0	4.0	
6	Student 7	3.0	4.0	4.0	1.0	1.0	1.0	4.0	
7	Student 8	3.0	4.0	2.0	2.0	3.0	2.0	4.0	
8	Student 9	3.0	4.0	4.0	2.0	3.0	1.0	4.0	
9	Student 10	4.0	4.0	3.0	3.0	2.0	2.0	4.0	
10	Student 11	4.0	5.0	4.0	2.0	3.0	3.0	5.0	
11	Student 12	3.0	3.0	3.0	1.0	1.0	2.0	5.0	
12	Student 13	2.0	3.0	3.0	1.0	1.0	2.0	4.0	
13	Student 14	3.0	3.0	4.0	2.0	2.0	2.0	4.0	
14	Student 15	4.0	4.0	5.0	3.0	3.0	1.0	4.0	
15	Student 16	3.0	3.0	4.0	1.0	1.0	1.0	4.0	
16	Student 17	4.0	4.0	3.0	1.0	1.0	1.0	5.0	
17	Student 18	3.0	4.0	3.0	2.0	1.0	3.0	4.0	
18	Student 19	3.0	3.0	4.0	1.0	2.0	1.0	4.0	
19	Student 20	1.0	5.0	3.0	1.0	1.0	2.0	1.0	
20	Student 21	3.0	4.0	3.0	2.0	3.0	1.0	4.0	
21	Student 22	4.0	5.0	5.0	1.0	1.0	2.0	4.0	
22	Student 23	2.0	2.0	2.0	1.0	1.0	4.0	4.0	
23	Student 24	3.0	5.0	5.0	2.0	1.0	1.0	4.0	

	CSL201	CSL202	DSL201
0	5.0	5.0	5.0
1	5.0	5.0	5.0
2	5.0	4.0	NaN
3	5.0	5.0	2.0

4	5.0	5.0	4.0
5	5.0	5.0	5.0
6	4.0	4.0	5.0
7	4.0	4.0	5.0
8	4.0	4.0	4.0
9	4.0	5.0	4.0
10	5.0	5.0	4.0
11	4.0	3.0	5.0
12	4.0	3.0	3.0
13	5.0	4.0	4.0
14	5.0	4.0	5.0
15	4.0	4.0	4.0
16	5.0	5.0	5.0
17	4.0	4.0	5.0
18	3.0	5.0	5.0
19	5.0	4.0	3.0
20	4.0	4.0	5.0
21	1.0	5.0	3.0
22	2.0	2.0	4.0
23	4.0	4.0	5.0

```
[18]: df.drop(["Student ID"],axis=1,inplace=True)
```

```
[19]: def introduce_missing_data(df, missing_percentage):
    df_missing = df.copy()
    num_missing = int(df.size * missing_percentage)
    missing_indices = np.random.choice(df.size, num_missing, replace=False)
    df_missing.values.ravel()[missing_indices] = np.nan
    return df_missing
```

```
[20]: def predict_missing_values(df):
    df_filled = df.copy()
    mse_scores = []

    for col in df.columns:
        if df[col].isnull().any():
            # Separate data into training and missing sets
            known_data = df_filled[df_filled[col].notnull()]
            unknown_data = df_filled[df_filled[col].isnull()]

            # Features and targets
            X_known = known_data.drop(columns=[col])
            y_known = known_data[col]

            X_unknown = unknown_data.drop(columns=[col])

            # Train a Linear Regression model
```



```

        model = LinearRegression()
        model.fit(X_known, y_known)

        # Predict missing values
        predicted_values = model.predict(X_unknown)
        df_filled.loc[unknown_data.index, col] = predicted_values

        # Evaluate using MSE (if possible)
        X_train, X_test, y_train, y_test = train_test_split(X_known, ↵
        ↪y_known, test_size=0.2, random_state=42)
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
        mse = mean_squared_error(y_test, y_pred)
        mse_scores.append(mse)

    return df_filled, mse_scores

```

```

[21]: def visualize_results(original, predicted):
        plt.figure(figsize=(10, 6))
        plt.scatter(original, predicted, alpha=0.6, edgecolors='k')
        plt.xlabel('Actual Values')
        plt.ylabel('Predicted Values')
        plt.title('Actual vs Predicted Values')
        plt.grid(True)
        plt.show()

```

```

[22]: missing_percentages = [0.2, 0.4, 0.6, 0.8]

for missing_percentage in missing_percentages:
    df_missing = introduce_missing_data(df, missing_percentage)
    df_filled, mse_scores = predict_missing_values(df_missing)
    visualize_results(df_missing, df_filled)

```

